Ashish Maurya

Location: Maharashtra, Mumbai, India

Ashishmaurya.com | LinkedIn | GitHub | LeetcodeEmail: ramashishmaurya738@gmail.com | Mobile: +91 9769295699

Summary

I am a motivated and analytical **Data Scientist** with strong skills in **Python, Machine Learning, SQL, and Data Analysis**. I have hands-on experience collecting, cleaning, and modeling data to identify patterns and develop ML solutions. I am proficient in **Pandas, NumPy, Scikit-learn, Matplotlib, and Seaborn**, and passionate about solving real-world business problems through Exploratory Data Analysis (EDA) and predictive modeling.

Skills

Programming: Python, R, SQL
Data Science & ML: Machine Learning, Exploratory Data Analysis (EDA), Data Visualization, Statistics
Libraries: Pandas, NumPy, Scikit-learn, Matplotlib, Seaborn, PyTorch, LangChain
Databases & Tracking: SQL, MongoDB, PostgreSQL, MLflow
Tools: Git, GitHub, Jupyter Notebook, Docker

Education

University of Mumbai — Maharashtra, Mumbai, India
Bachelor of Data Science — *Aug 2022 – Jun 2025*

Projects

NeuralDocs AI | Reactjs, FastAPI, LangChain, RAG, Qdrant, Docker, WebSockets, Redis — [View on GitHub]

- Built a full-stack RAG assistant (**React, FastAPI**) enabling users to upload and query complex PDFs in natural language.
- Integrated real-time **WebSocket streaming with LangChain** to deliver instant, word-by-word AI responses with zero latency.
- Designed an asynchronous ingestion pipeline **using Celery and Redis for seamless, non-blocking background processing** of large files.
- Optimized retrieval performance and cost by **leveraging Qdrant vector database**, context-aware querying, and LLM caching.

JobSalaryPrediction | Python, Random Forest, XGBoost, MLflow, Docker, Render, Streamlit — [View on GitHub]

- Built a **Job Salary Prediction system** using **Machine Learning** models (Random Forest, XGBoost) trained on real-world job listings to predict salaries based on role, experience, location, and skills with high accuracy.
- Integrated **MLflow** for complete experiment tracking — logged model parameters, metrics, and artifacts across multiple runs to compare and select the best performing model for production deployment.
- Containerized the full ML pipeline using **Docker** and deployed on **Render** for public access, enabling real-time salary predictions through an interactive **Streamlit** web application.

Email-spam-detection | Python, NLP, TF-IDF, Naive Bayes, Logistic Regression — [View on GitHub]

- Built an **Email Spam Detection system** using **NLP techniques** — preprocessed email text with tokenization, stop-word removal and **TF-IDF vectorization** to classify emails as spam or not spam.
- Trained and compared **Naive Bayes and Logistic Regression** models on real email datasets, achieving **95% accuracy** with a clean **Streamlit** interface for live email classification.

Certifications

- Fundamental Tools of Data Wrangling
- Complete Data Science, Machine Learning, DL, NLP
- Structure Query Language (SQL)
- Data Structure And algorithm Using Python